

REPORT

-Krishna Shukla(2022254)

Theoretical Questions(Questions 1 and 2)

CV Axon - 2

As Rot is about y axis by angle θ

$$R_y(\theta) = \begin{pmatrix} \cos\theta & 0 & \sin\theta \\ 0 & 1 & 0 \\ -\sin\theta & 0 & \cos\theta \end{pmatrix} \quad R_z(-\pi) = \begin{pmatrix} \sqrt{3}/2 & 0 & -1/2 \\ 0 & 1 & 0 \\ \sqrt{3}/2 & 0 & 1/2 \end{pmatrix}$$
$$R_x(\theta) = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \cos\theta & -\sin\theta \\ 0 & \sin\theta & \cos\theta \end{pmatrix} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & \sqrt{2}/2 & -\sqrt{2}/2 \\ 0 & \sqrt{2}/2 & \sqrt{2}/2 \end{pmatrix}$$

Reflection across XY plane implies $R_{xy} = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix}$

Part - 2

$$C = e_{xx} R_x(\pi/4) R_y(\pi/4) R_z(-\pi)$$
$$\cdot C = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} \sqrt{2}/2 & 0 & -1/2 \\ -1/2 & \sqrt{2}/2 & \sqrt{2}/2 \\ 1/2 & \sqrt{2}/2 & -\sqrt{2}/2 \end{pmatrix} \begin{pmatrix} \sqrt{3}/2 & 0 & 1/2 \\ 0 & 1 & 0 \\ 1/2 & 0 & \sqrt{2}/2 \end{pmatrix}$$
$$\therefore C = \begin{pmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} \sqrt{3}/2 & 0 & -1/2 \\ -\frac{1}{2\sqrt{2}} & \frac{\sqrt{2}}{2\sqrt{2}} & \frac{\sqrt{2}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{\sqrt{2}}{2\sqrt{2}} & -\frac{\sqrt{2}}{2\sqrt{2}} \end{pmatrix} = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & \frac{1}{2} \\ \frac{1}{2\sqrt{2}} & -\frac{1}{2} & \frac{\sqrt{2}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{1}{2} & -\frac{\sqrt{2}}{2\sqrt{2}} \end{pmatrix}$$

CD

$v_{new} = Cv + t$ where $v = \begin{pmatrix} 3 \\ 4 \\ 4 \end{pmatrix}$

$$v_{new} = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & \frac{1}{2} \\ \frac{1}{2\sqrt{2}} & -\frac{1}{2} & \frac{\sqrt{2}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{1}{2} & -\frac{\sqrt{2}}{2\sqrt{2}} \end{pmatrix} \begin{pmatrix} 3 \\ 4 \\ 4 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ -2 \end{pmatrix}$$
$$v_{new} = \boxed{\begin{pmatrix} \frac{3\sqrt{3}-2}{2} \\ \frac{\sqrt{3}+4\sqrt{2}}{2\sqrt{2}} \\ \frac{1+4\sqrt{3}-4\sqrt{2}}{2\sqrt{2}} \end{pmatrix}}$$

$$u = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} \rightarrow u_{\text{new}} = \begin{pmatrix} u_1 \\ u_2 \\ u_3 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix} + \begin{pmatrix} 1 \\ 0 \\ -2 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ -2 \end{pmatrix} = \text{origin}$$

Combined Rotation = $R_x(\frac{\pi}{4}) R_y(-\frac{\pi}{6}) = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ \frac{1}{2}\sqrt{2} & \frac{1}{2} & -\frac{\sqrt{3}}{2}\sqrt{2} \\ \frac{1}{2}\sqrt{2} & \frac{1}{2} & \frac{\sqrt{3}}{2}\sqrt{2} \end{pmatrix}$

$$\theta = \arccos \left(\frac{\text{tr}(R) - 1}{2} \right)$$

$$\theta = \arccos \left(\frac{\sqrt{3} + 2 + \sqrt{3} - 2\sqrt{2}}{4\sqrt{2}} \right) \approx 55^\circ$$

$$u = \frac{1}{2\sin\theta} \begin{pmatrix} C_{x3} - C_{x2} \\ C_{y3} - C_{y2} \\ C_{z3} - C_{z2} \end{pmatrix} = u = \begin{pmatrix} (\frac{1}{2} + \frac{\sqrt{3}}{2})/2\sin(55^\circ) \\ (\frac{1}{2} - \frac{1}{2}\sqrt{2})/2\sin(55^\circ) \\ (\frac{1}{2} - 0)/2\sin(55^\circ) \end{pmatrix}$$

$$u = \begin{pmatrix} 0.81 \\ -0.53 \\ -0.21 \end{pmatrix} \rightarrow \text{Axis of Rotn}$$

$$R = I + \sin\theta [u]_x + (1-\cos\theta) [u]_x^2$$

θ = Angle of Rotn (Up to 3rd symm matrix \rightarrow)

$$R = I + 0.805 \begin{pmatrix} 0 & 0.21 & -0.53 \\ -0.21 & 0 & -0.81 \\ 0.53 & 0.81 & 0 \end{pmatrix} + 0.407 \begin{pmatrix} 0 & -u_2 & u_1 \\ u_2 & 0 & -u_3 \\ -u_1 & u_3 & 0 \end{pmatrix}$$

$$= \begin{pmatrix} 0.805 & 0 & -0.50 \\ -0.35 & 0.70 & -0.61 \\ -0.33 & 0.70 & 0.61 \end{pmatrix} \approx \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2}\sqrt{2} \\ \frac{1}{2}\sqrt{2} & \frac{1}{2} & -\frac{\sqrt{3}}{2}\sqrt{2} \\ \frac{1}{2}\sqrt{2} & \frac{1}{2} & \frac{\sqrt{3}}{2}\sqrt{2} \end{pmatrix}$$

Hence proved

Hence Rodrigues and Rotation give the same answer:

Q.2 [Theory]

$$\mathbf{n} = \mathbf{K}(\mathbf{R}\mathbf{I} + \mathbf{t})\mathbf{x}, \quad (\mathbf{t} = \mathbf{0} \text{ as pure rotation R})$$

$$\text{Camera 1} \rightarrow \mathbf{n}_1 = \mathbf{K}_1(\mathbf{R}_1\mathbf{I}_1) \mathbf{x}$$

$$\text{Cam 2} \rightarrow \mathbf{n}_2 = \mathbf{K}_2(\mathbf{R}_2\mathbf{I}_2) \mathbf{x}$$

$$\mathbf{n}_1 = \mathbf{K}_1 \mathbf{x} \Rightarrow \mathbf{x}_1 = \mathbf{K}_1^{-1} \mathbf{n}_1$$

$$\mathbf{n}_2 = \mathbf{K}_2 \mathbf{R}_2 \mathbf{x}$$

$$\mathbf{n}_2 = \mathbf{K}_2 \mathbf{R}_2 (\mathbf{K}_1^{-1} \mathbf{n}_1)$$

$$\mathbf{n}_2 = (\mathbf{K}_2 \mathbf{R}_2 \mathbf{K}_1^{-1}) \mathbf{n}_1 \rightsquigarrow \mathbf{n}_2 = \mathbf{H}^T \mathbf{n}_1$$

$$\mathbf{n}_1 = (\mathbf{H}^T)^{-1} \mathbf{n}_2$$

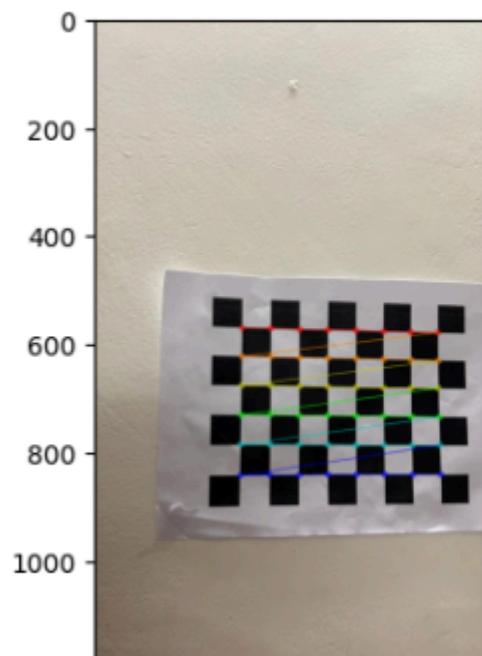
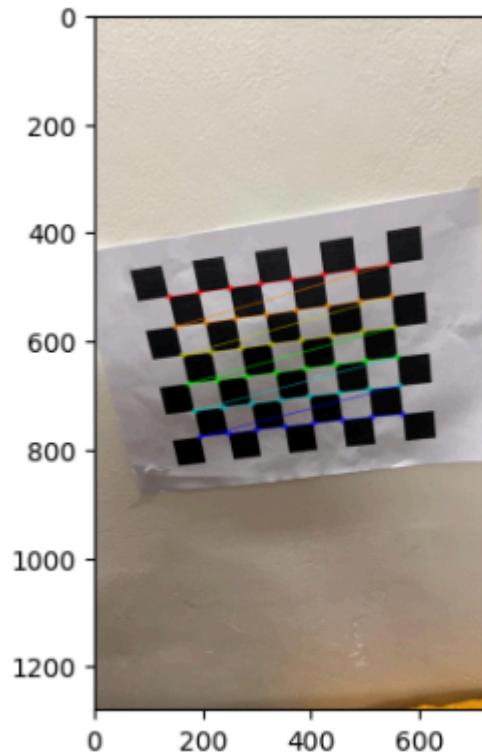
$$= (\mathbf{K}_1^{-1} \mathbf{R}_2^{-1} \mathbf{K}_2^{-1}) \mathbf{n}_2$$

As \mathbf{K}_1 & \mathbf{K}_2 are camera intrinsic param, they are invertible
 \mathbf{R} is a pure rot'g matrix (orthogonal $\mathbf{R}^T = \mathbf{R}^{-1}$), thus identity

Coding Question(Questions 3,4)

Camera Calibration:

The image visualization for all image 2 examples are shown below



1)

```
Camera matrix :  
  
[[956.63719897  0.          369.05485204]  
 [ 0.          957.55141318  651.41419982]  
 [ 0.          0.          1.          ]]  
fx = 956.637198973105  
fy = 957.5514131759558  
cx = 369.0548520367341  
cy = 651.414199816039  
s = 0.0
```

2)

```
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/1.jpeg:  
Rotation Matrix:  
[[ 0.91736941 -0.19600125 -0.34643452]  
 [-0.01803989  0.84899005 -0.5281008 ]  
 [ 0.39762788  0.49071316  0.77530166]]  
Translation Vector:  
[[-3.93829473]  
 [-3.61716296]  
 [14.65911195]]  
  
[[ 0.99699847  0.01387206 -0.0761683 ]  
 [-0.02511535  0.9885627  -0.14870443]  
 [ 0.07323431  0.15017108  0.9859439 ]]  


---

  
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/2.jpeg:  
Rotation Matrix:  
[[ 0.91736941 -0.19600125 -0.34643452]  
 [-0.01803989  0.84899005 -0.5281008 ]  
 [ 0.39762788  0.49071316  0.77530166]]  
Translation Vector:  
[[-1.49771589]  
 [-0.657919 ]  
 [13.91817088]]
```

In notebook printed for all images

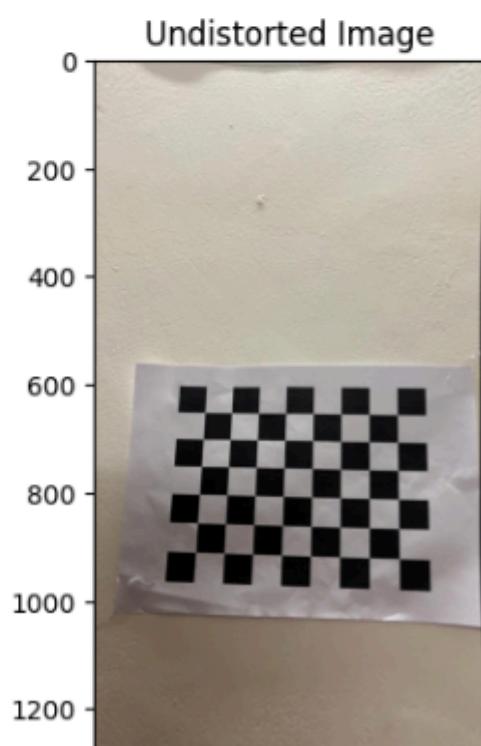
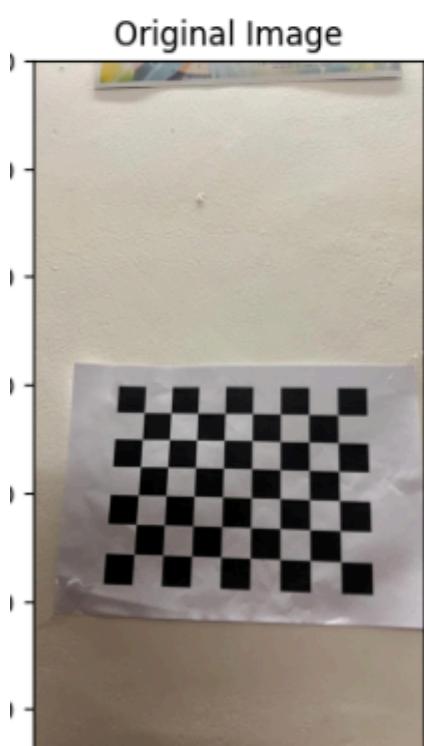
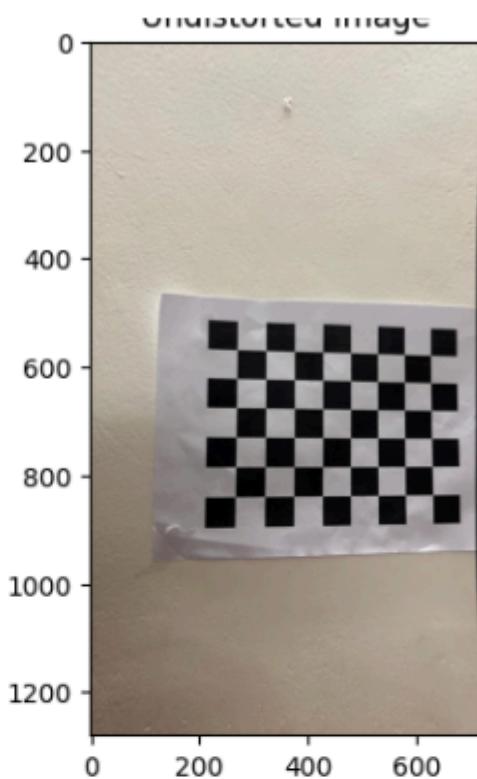
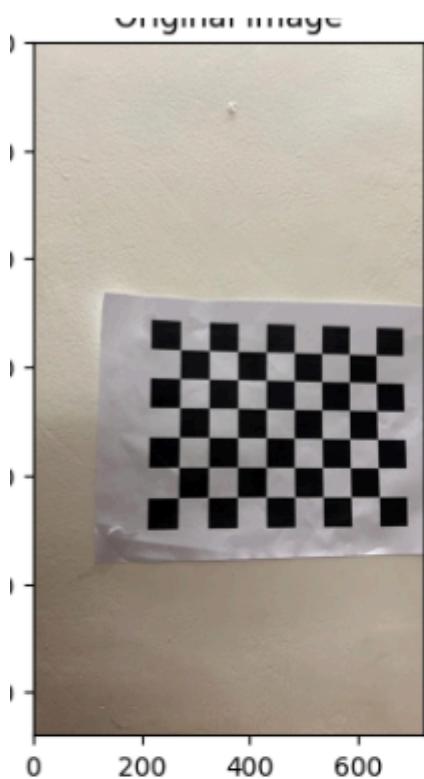
3)

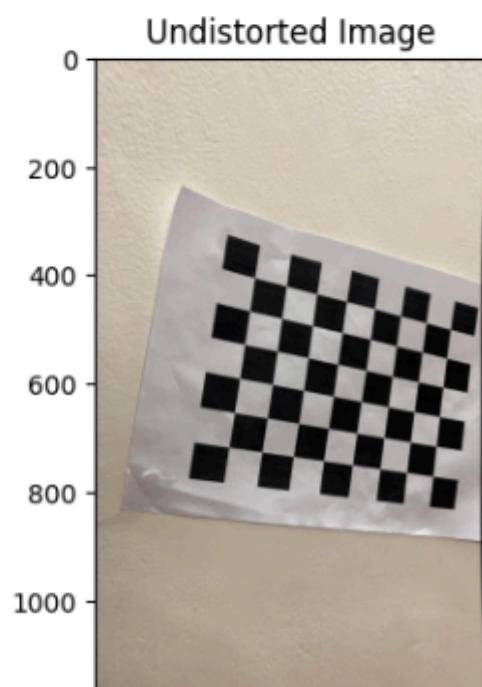
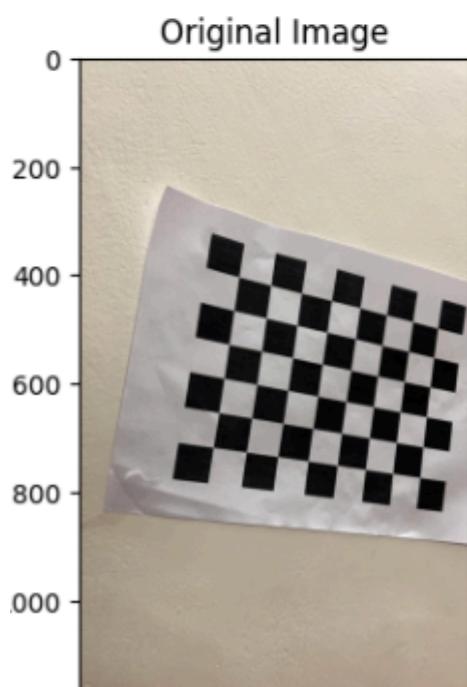
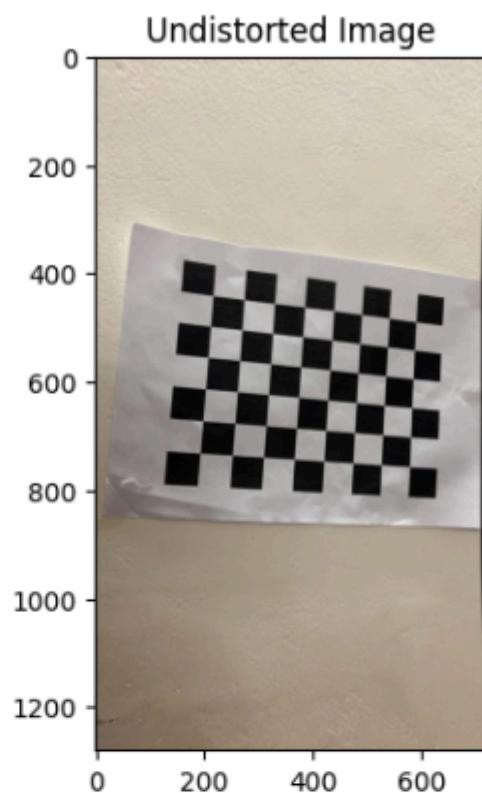
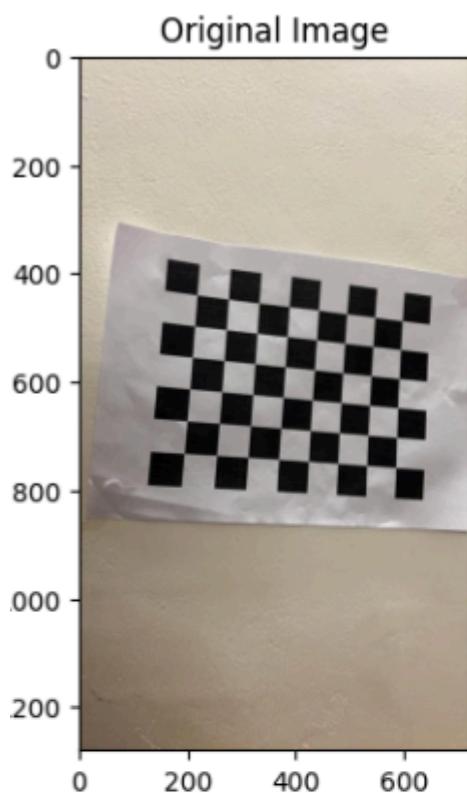
```
Distortion coefficients: [[ 0.19763718 -0.70686979  0.00336936  0.00698242  0.04879612]]
```

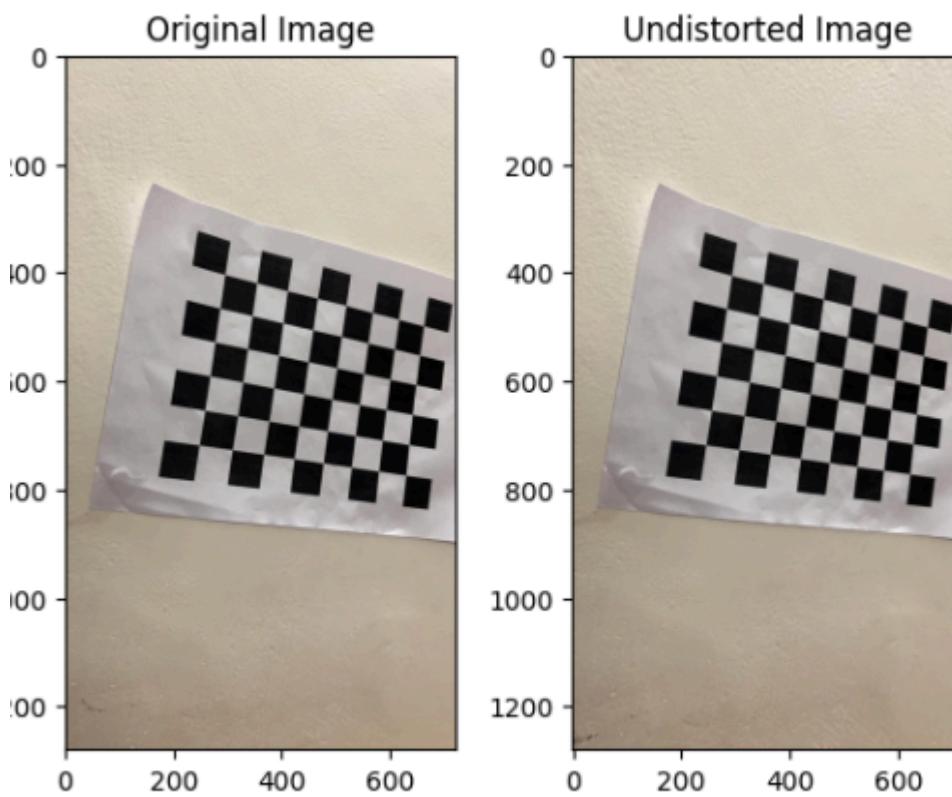
Where 0th value, 1st value, 4th value are k1,k2,k3 that is radial distortion coefficient whereas 2nd and 3rd value are p1 and p2 which is tangential distortion coefficient

k1, k2, k3 = 0.19763718, -0.70686979, 0.04879612
p1, p2 = 0.00336936, 0.00698242

[Radial]
[Tangential]



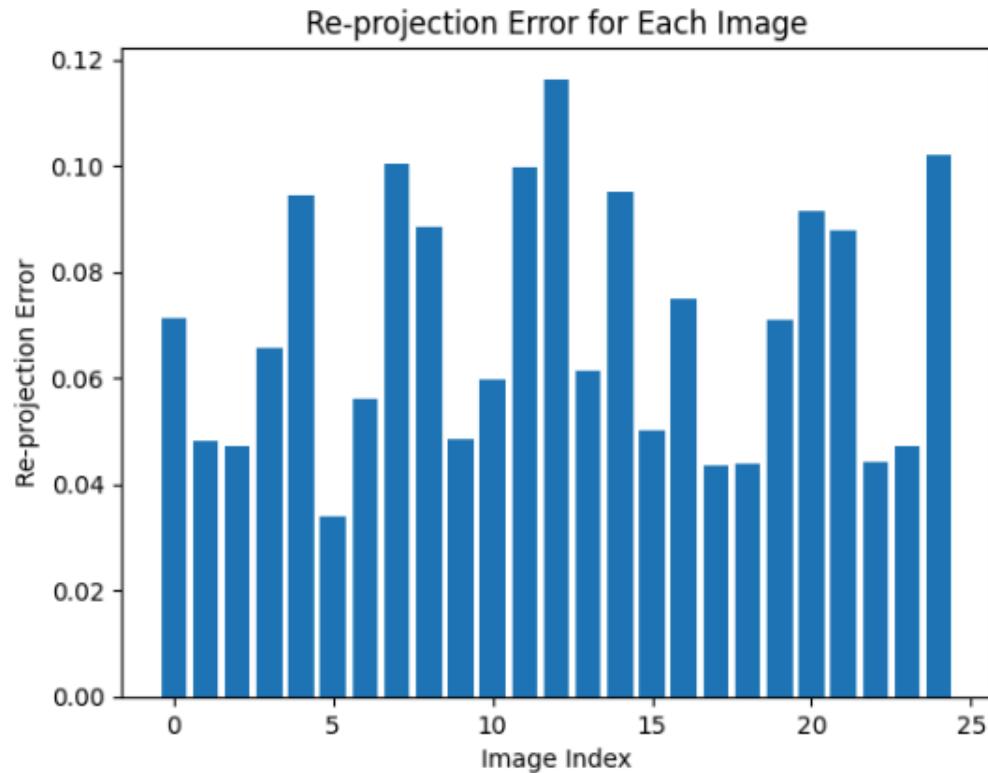




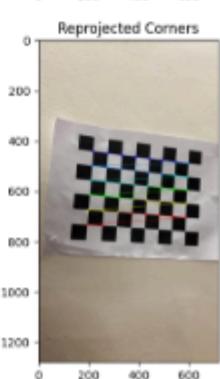
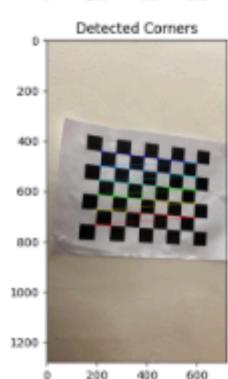
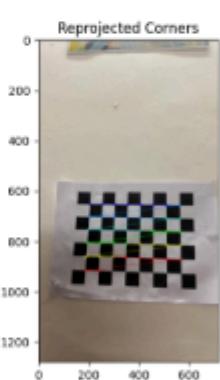
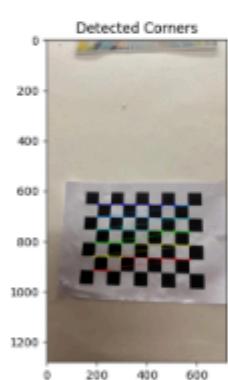
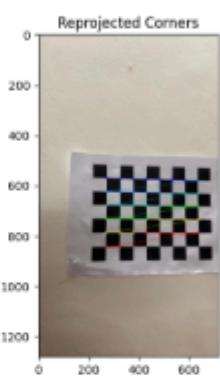
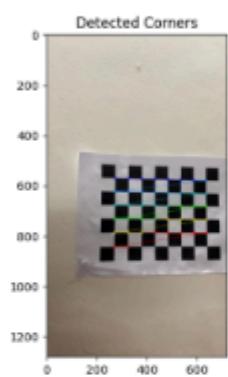
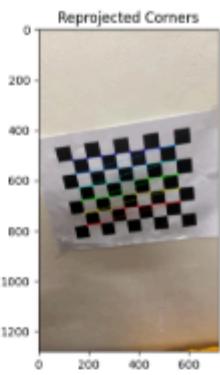
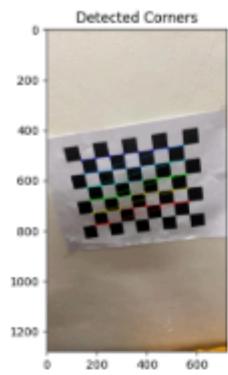
Comment-

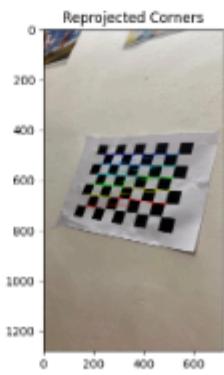
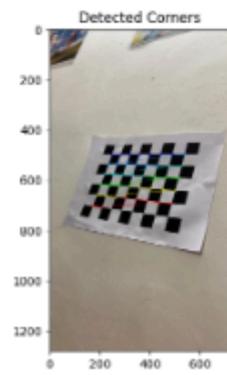
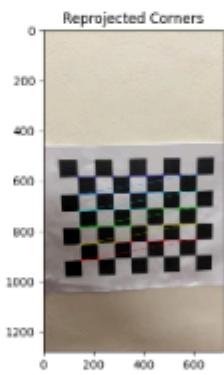
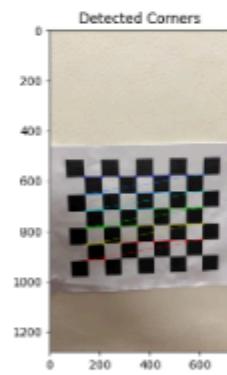
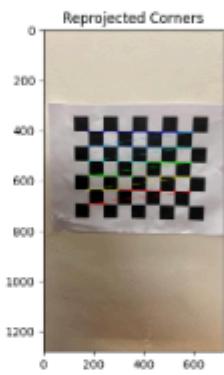
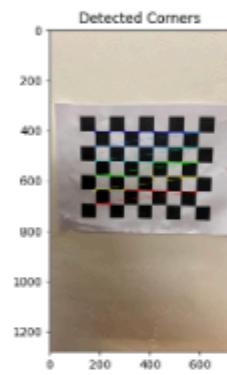
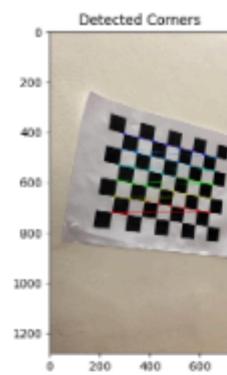
After applying the radial distortion coefficients to undistort the images, it is observed that the previously curved straight lines, especially those located at the corners, have been corrected to appear straight. The distortion, which caused the lines to bow outward or inward, has been effectively compensated, resulting in a more geometrically accurate representation of the scene. This correction is particularly noticeable in the corners, where the distortion was most pronounced. The undistorted images now exhibit improved clarity and accuracy, highlighting the effectiveness of the radial distortion coefficients in restoring the true geometry of the captured scenes.

4)



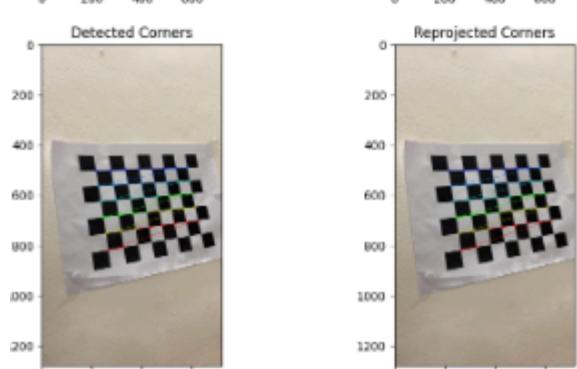
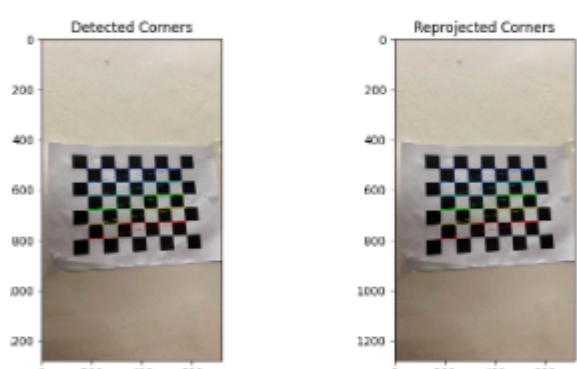
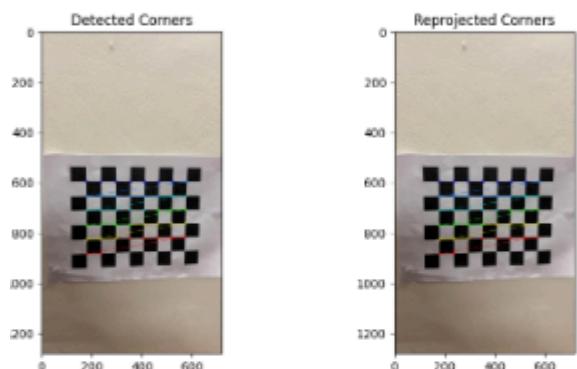
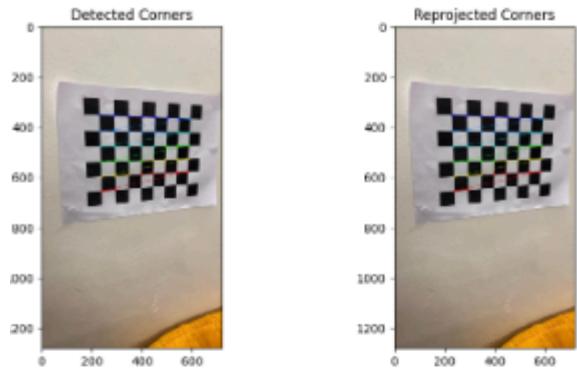
The reprojection error is computed by measuring the difference between the observed image points (i.e., the actual 2D points detected in the image) and the projected points obtained by projecting the corresponding 3D points back onto the image plane using the estimated camera calibration parameters. This involves first transforming the 3D world coordinates into the camera coordinate system using the extrinsic parameters (rotation and translation vectors), and then projecting these points onto the image plane using the intrinsic parameters (focal length, principal point, and distortion coefficients). The reprojection error is then calculated as the Euclidean distance between the observed 2D points and the projected 2D points for each image point. This error is typically averaged over all points and images to obtain the mean reprojection error, which provides an indication of the accuracy of the camera calibration process. A lower reprojection error signifies a more accurate calibration.

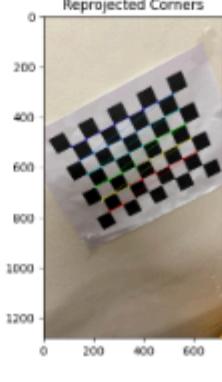
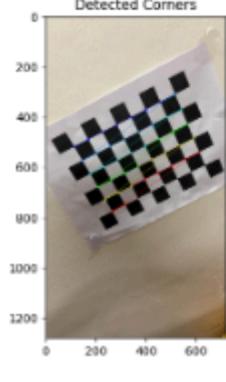
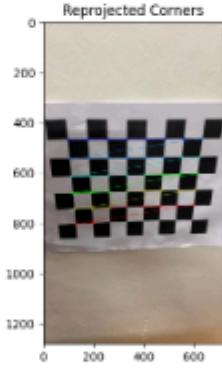
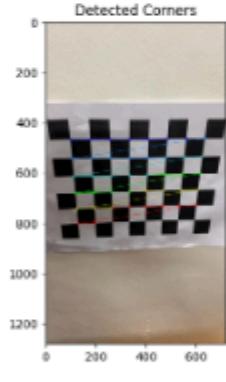
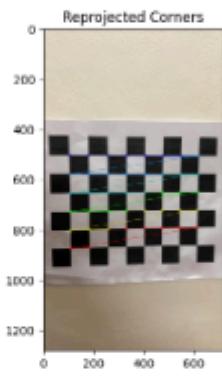
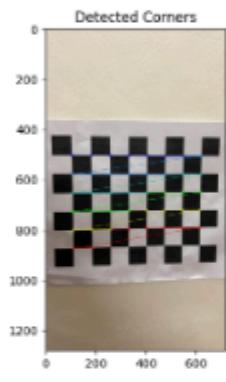
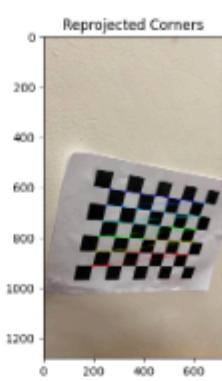
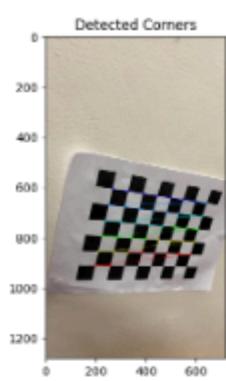


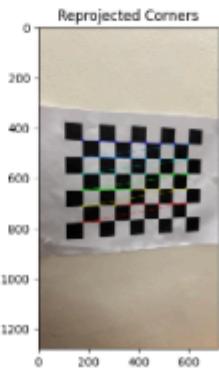
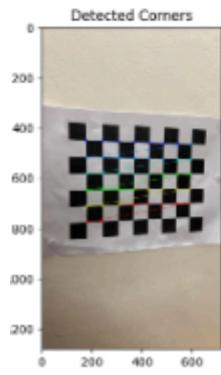
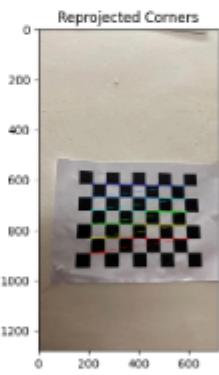
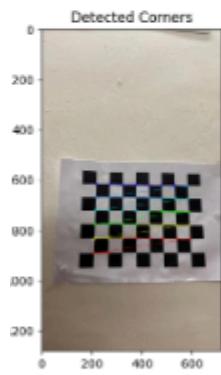
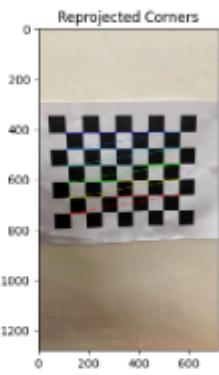
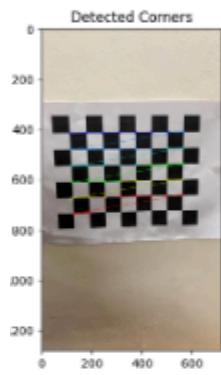
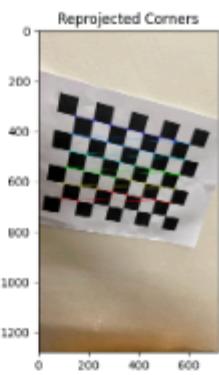
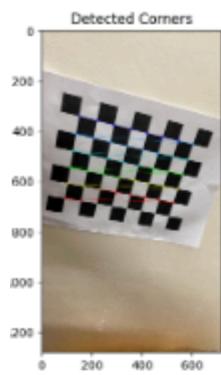


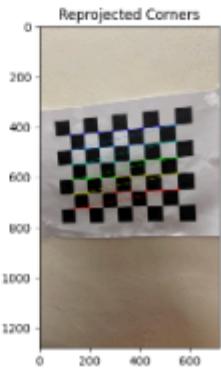
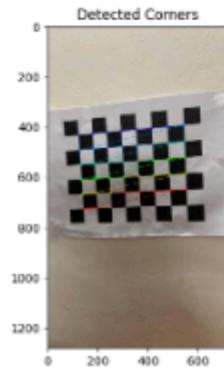
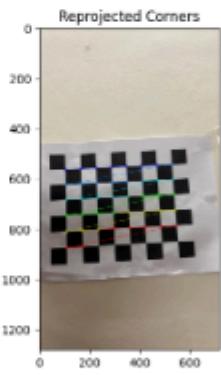
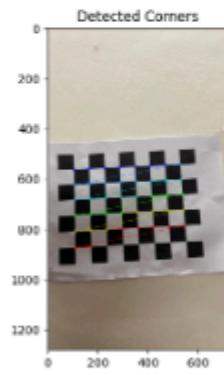
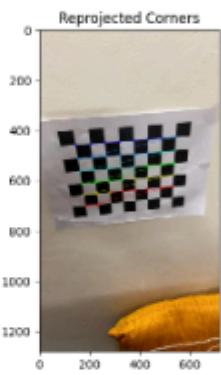
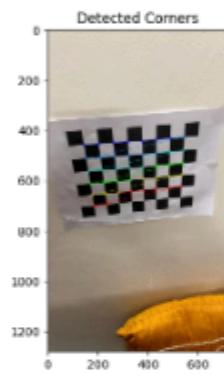
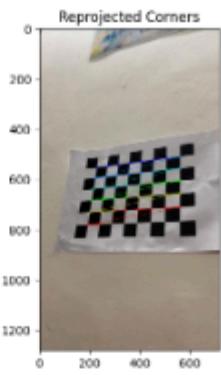
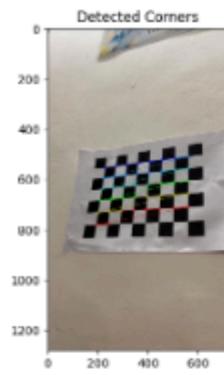
Distorted Corners

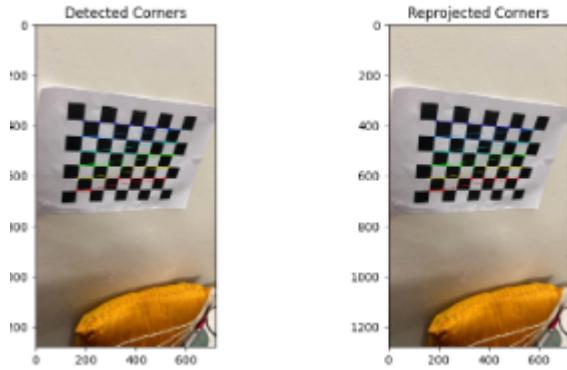
Reprojected Corners











6)

```

/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/14.jpeg:
Plane Normal: [ 0.04521854 -0.32517343  0.94457267]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/18.jpeg:
Plane Normal: [-0.14608872  0.07386587  0.98651115]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/16.jpeg:
Plane Normal: [ 0.07201336 0.21723429 0.97345947]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/25.jpeg:
Plane Normal: [-0.22325172  0.06830378  0.97236478]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/22.jpeg:
Plane Normal: [-0.36034965  0.01814488  0.93264089 ]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/24.jpeg:
Plane Normal: [ 0.0356678 -0.11054531  0.99323086]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/6.jpeg:
Plane Normal: [-0.08327302 -0.12752644  0.98833325]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/9.jpeg:
Plane Normal: [ 0.45456232 0.48009243 0.79580095]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/11.jpeg:
Plane Normal: [-0.39826967 -0.1871547  0.89797238]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/20.jpeg:
Plane Normal: [-0.13551124 -0.10599322  0.98508991]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/4.jpeg:
Plane Normal: [-0.16581683 0.16012457 0.97306983]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/5.jpeg:
Plane Normal: [-0.39737416 0.19880494 0.89760883]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/2.jpeg:
Plane Normal: [-0.38526818 -0.30920821 0.86945886]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/19.jpeg:
Plane Normal: [-0.1077335 -0.08642536 0.99041615]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/7.jpeg:
Plane Normal: [-0.18314688 -0.40167845 0.909957 ]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/21.jpeg:
Plane Normal: [-0.09763584 -0.27445986 0.9566293 ]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/23.jpeg:
Plane Normal: [-0.16459244 -0.36280349 0.91721478]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/1.jpeg:
Plane Normal: [-0.0761683 -0.14870443 0.9859439 ]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/17.jpeg:
Plane Normal: [-0.09543046 0.19721867 0.97570537]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/3.jpeg:
Plane Normal: [-0.27149228 -0.09615281 0.95765188]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/13.jpeg:
Plane Normal: [ 0.35820577 0.39861417 0.8442721 ]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/15.jpeg:
Plane Normal: [-0.12778898 -0.50726445 0.85226331]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/10.jpeg:
Plane Normal: [-0.0228698 0.12383015 0.99203985]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/8.jpeg:
Plane Normal: [ 0.18876585 -0.06635959 0.97977745]
/kaggle/input/chess-board/chessboard_dataset-20250402T180104Z-001/chessboard_dataset/12.jpeg:
Plane Normal: [-0.34643452 -0.5281008 0.77530166]

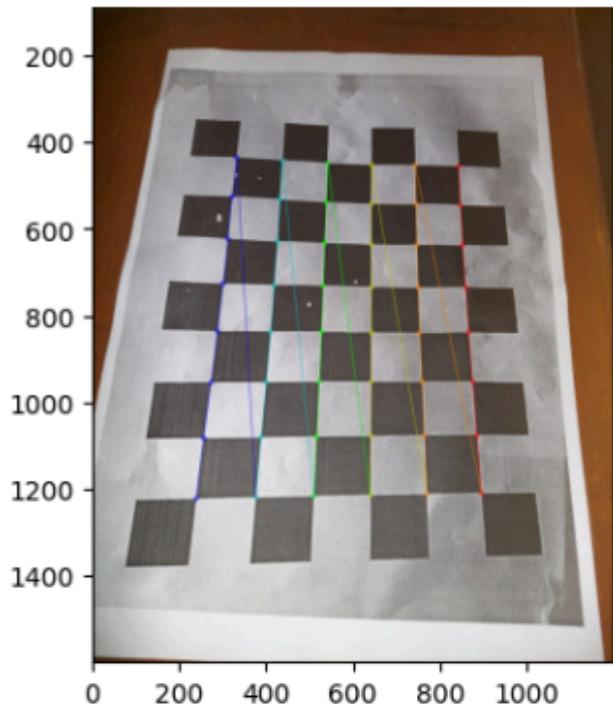
```

A set of small, semi-transparent navigation icons typically used in Jupyter notebooks for navigating between cells and sections.

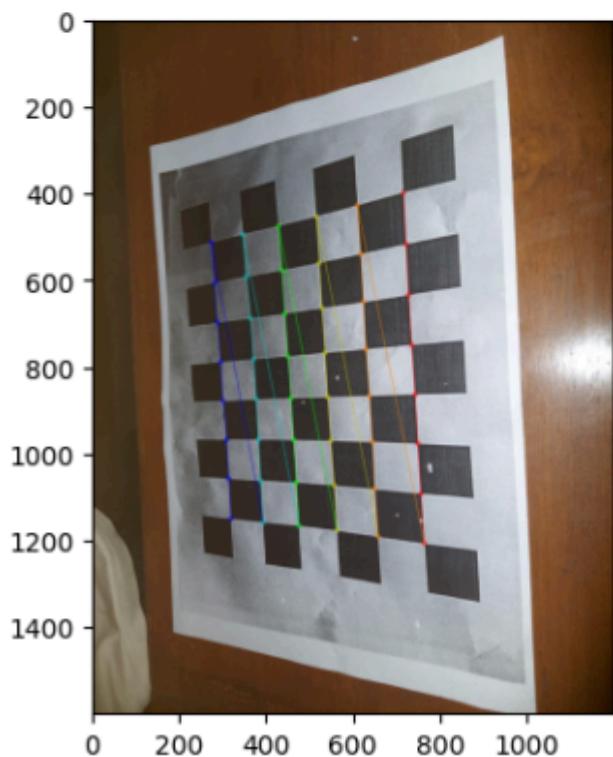
For My Dataset (Repeating)

Link to Dataset - [Dataset2](#)

All visualization done like this



`/kaggle/input/chess-board/Dataset2/Dataset2/IMG-2025040`



1)

```
Camera matrix :
```

```
[[1.21639845e+03 0.00000000e+00 7.61817209e+02]
 [0.00000000e+00 1.19963824e+03 8.67839611e+02]
 [0.00000000e+00 0.00000000e+00 1.00000000e+00]]
fx = 1216.3984514424942
fy = 1199.638241083678
cx = 761.8172089419866
cy = 867.8396114071278
s = 0.0
```

2)

```
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0062.jpg:
Rotation Matrix:
[[ 0.99521202 -0.01594962 -0.09642944]
 [ 0.05249355  0.91944992  0.38968739]
 [ 0.08244668 -0.3928835   0.91588487]]
Translation Vector:
[[-2.92962216]
 [-6.63911901]
 [10.59572754]]

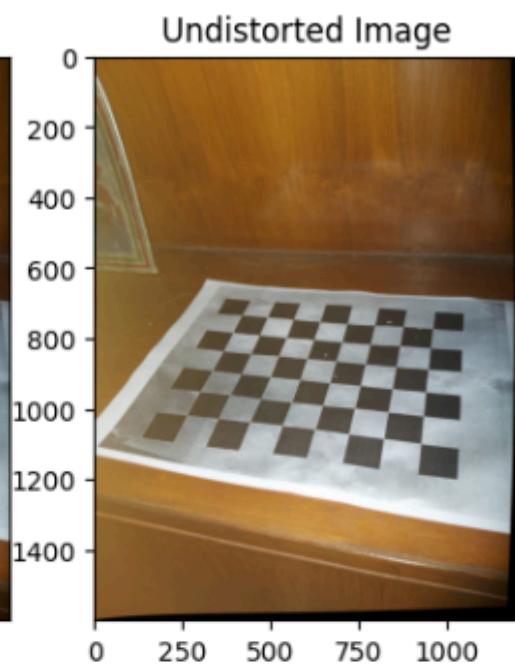
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0059.jpg:
Rotation Matrix:
[[ 0.99940419  0.00829917 -0.03350204]
 [ 0.0042843   0.93332392  0.35900989]
 [ 0.03424774 -0.35893953  0.93273228]]
Translation Vector:
[[-2.8147957 ]
 [-6.53792622]
 [11.30729187]]
```

In notebook printed for all images

3)

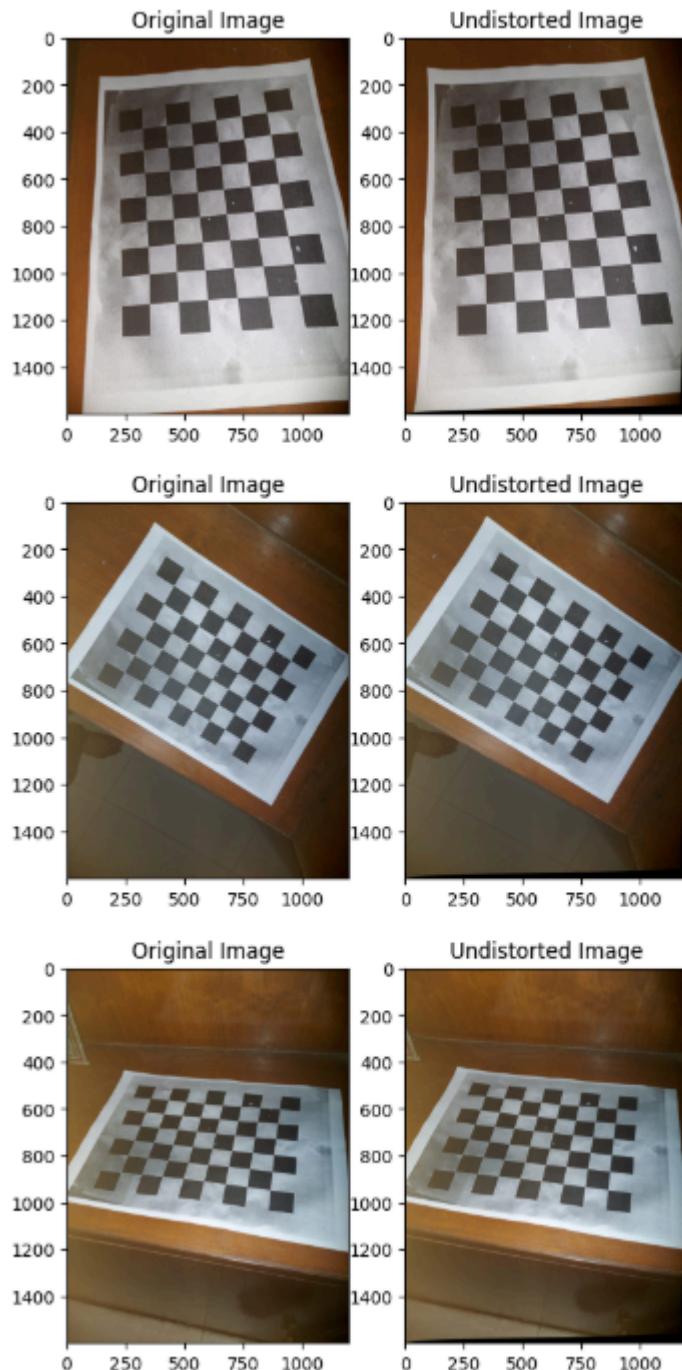
```
Distortion coefficients: [[-0.006025    0.18945388   0.01368964   0.02305314 -0.24908978]]
```

Parameter reasoning already given in above part



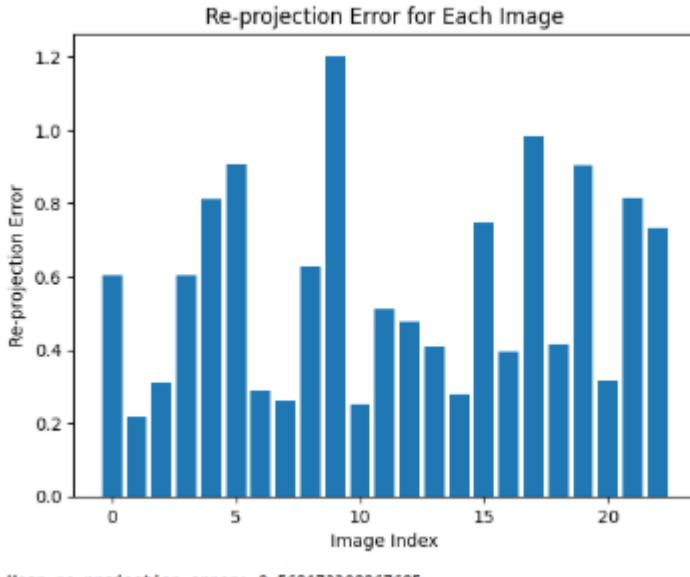
Original Image

Undistorted Image



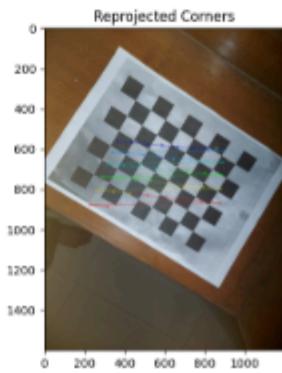
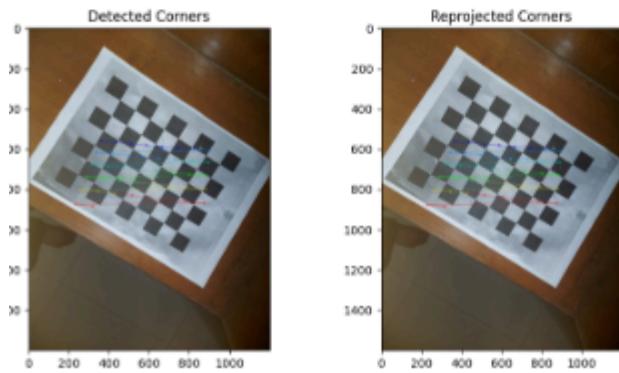
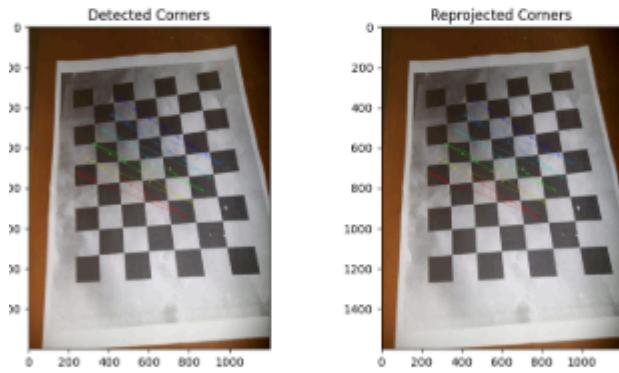
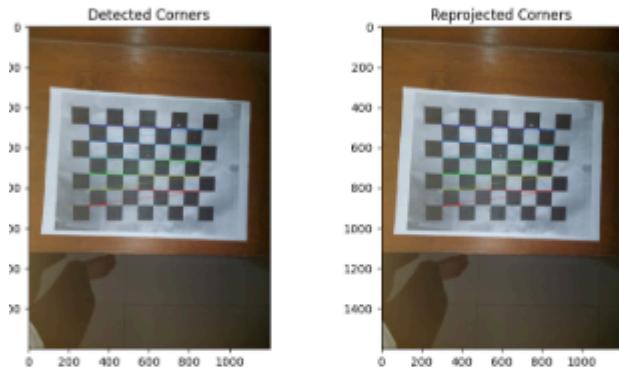
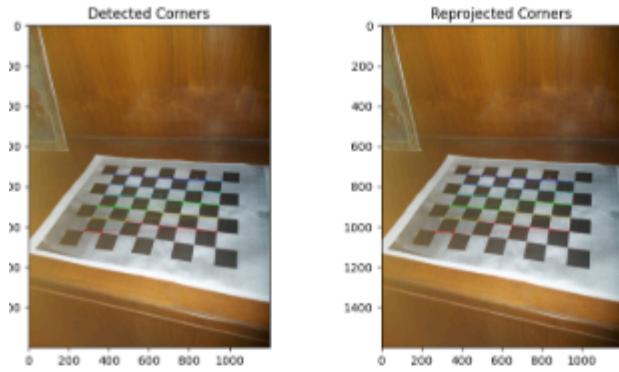
Comment-

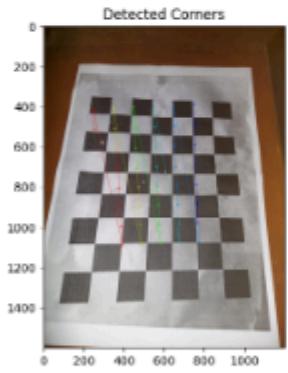
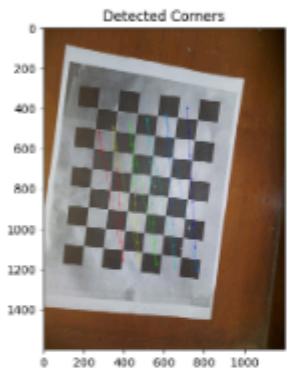
After applying the radial distortion coefficients to undistort the images, it is observed that the previously curved straight lines, especially those located at the corners, have been corrected to appear straight. The distortion, which caused the lines to bow outward or inward, has been effectively compensated, resulting in a more geometrically accurate representation of the scene. This correction is particularly noticeable in the corners, where the distortion was most pronounced. The undistorted images now exhibit improved clarity and accuracy, highlighting the effectiveness of the radial distortion coefficients in restoring the true geometry of the captured scenes.

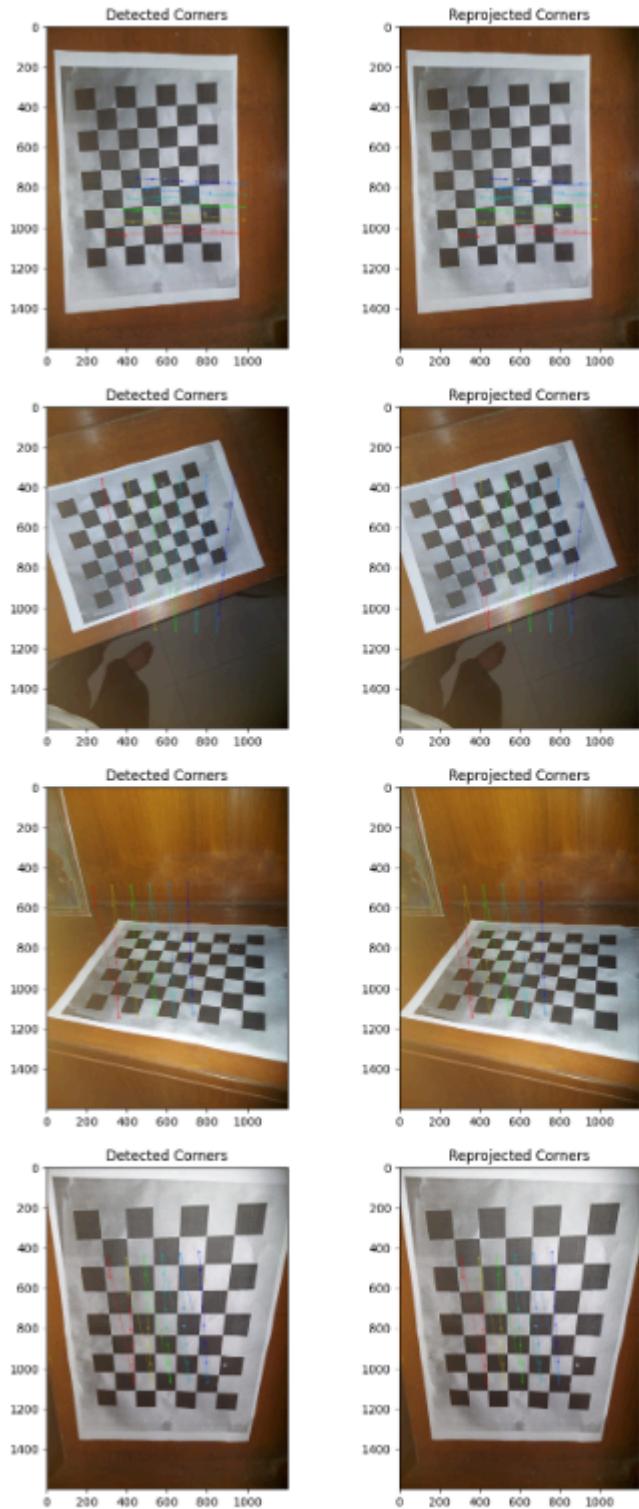


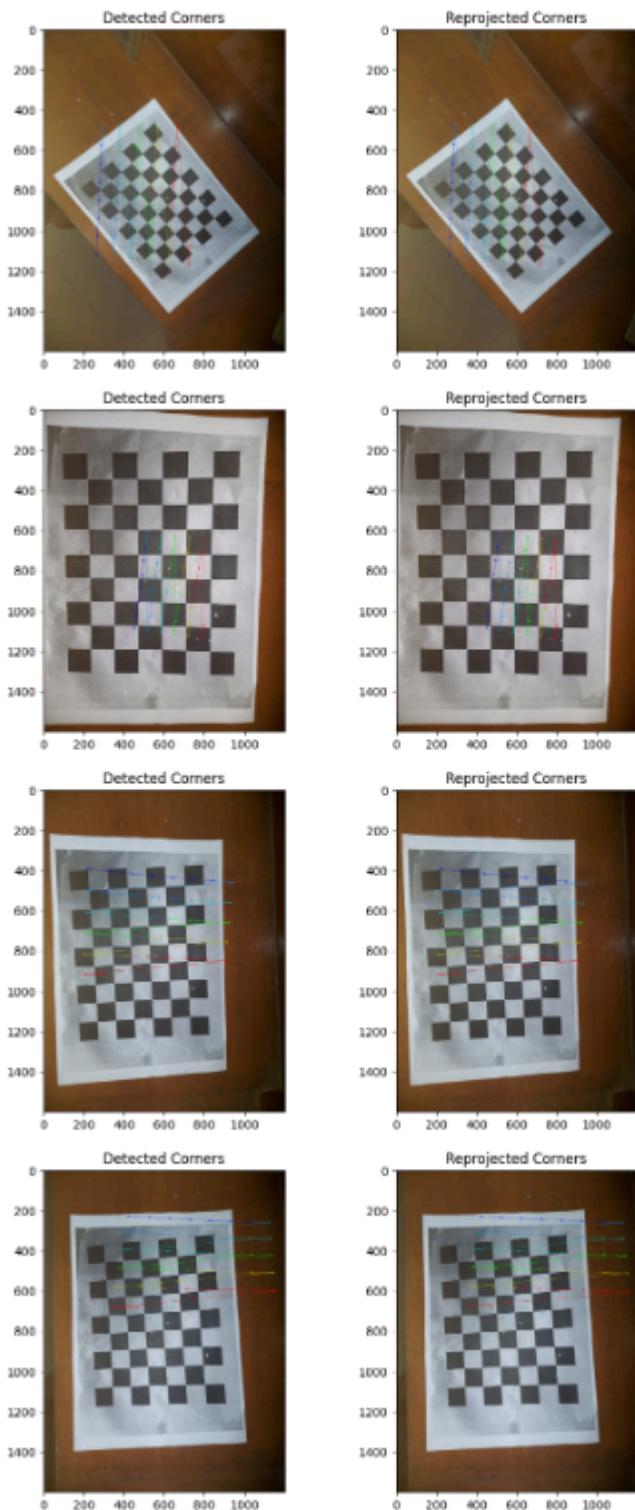
The reprojection error is computed by measuring the difference between the observed image points (i.e., the actual 2D points detected in the image) and the projected points obtained by projecting the corresponding 3D points back onto the image plane using the estimated camera calibration parameters. This involves first transforming the 3D world coordinates into the camera coordinate system using the extrinsic parameters (rotation and translation vectors), and then projecting these points onto the image plane using the intrinsic parameters (focal length, principal point, and distortion coefficients). The reprojection error is then calculated as the Euclidean distance between the observed 2D points and the projected 2D points for each image point. This error is typically averaged over all points and images to obtain the mean reprojection error, which provides an indication of the accuracy of the camera calibration process. A lower reprojection error signifies a more accurate calibration.

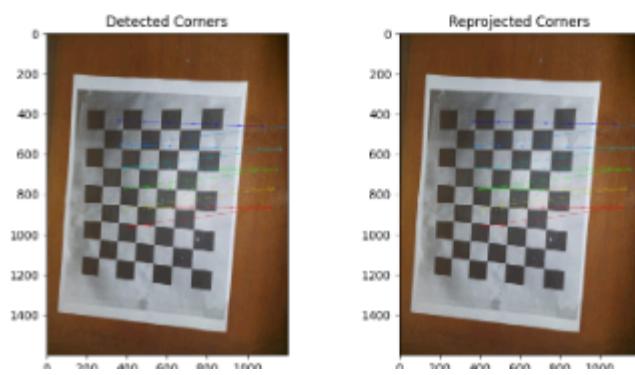
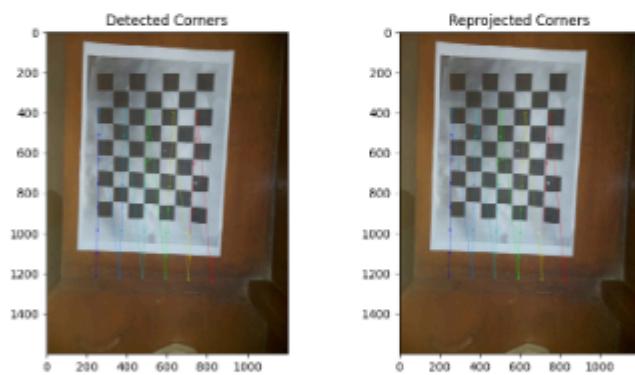
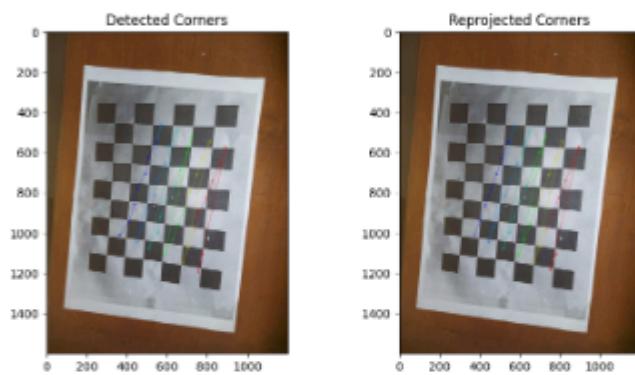
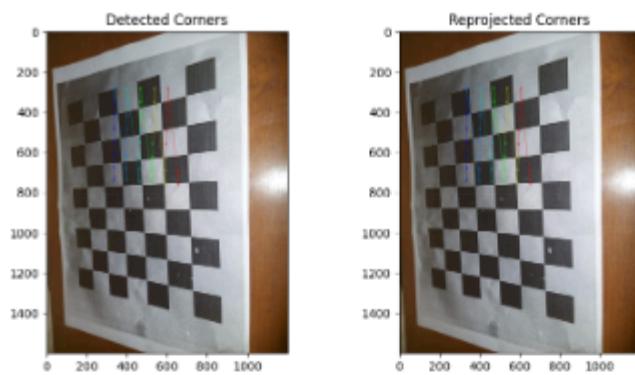
5)

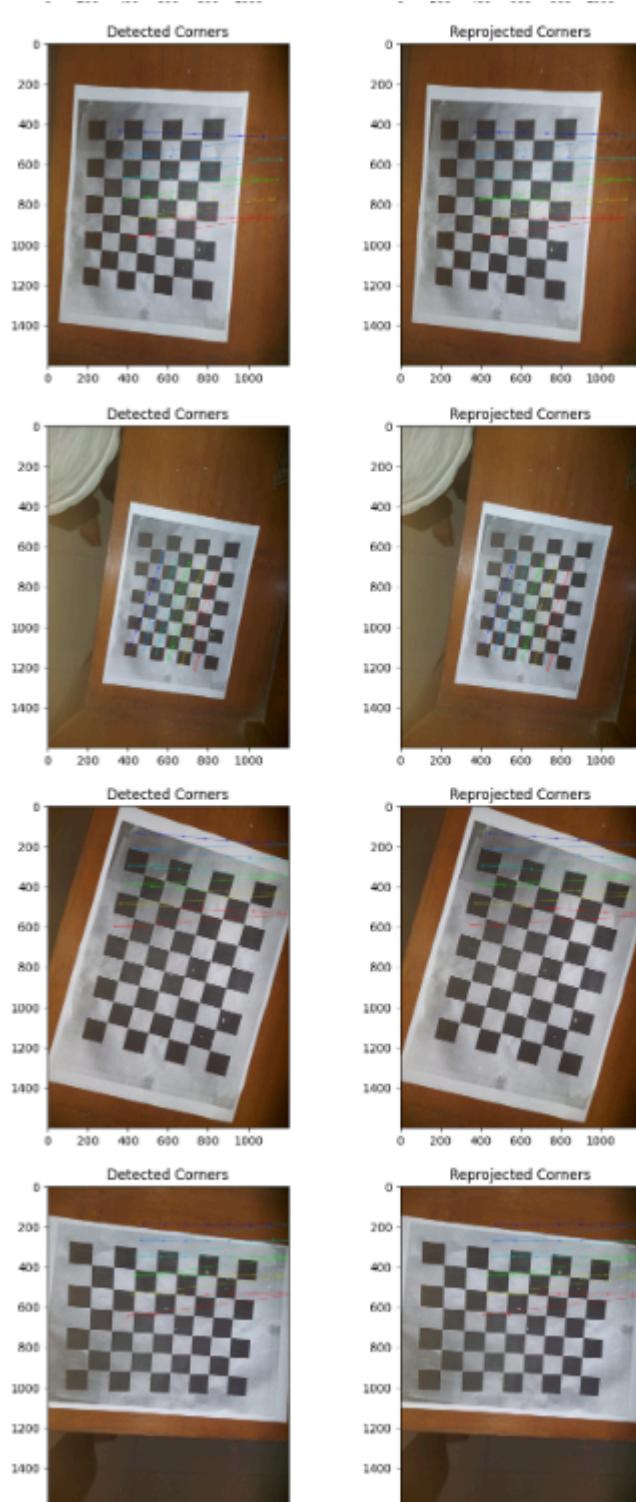












6)

```
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0047.jpg:  
Plane Normal: [ 0.0263037  0.72255562  0.6908122 ]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0067.jpg:  
Plane Normal: [-0.16616492 -0.16581171  0.97205745]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0042.jpg:  
Plane Normal: [-0.11038055 -0.19722716  0.97412401]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0061.jpg:  
Plane Normal: [-0.02630916  0.53553288  0.84410498]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0046.jpg:  
Plane Normal: [-0.13324063  0.29406317  0.94645327]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0070.jpg:  
Plane Normal: [ 0.3695629 -0.08272866  0.92551566]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0043.jpg:  
Plane Normal: [-0.02164137 -0.12354672  0.99210275]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0050.jpg:  
Plane Normal: [-0.23631129 -0.09835753  0.96668649]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0073.jpg:  
Plane Normal: [ 0.03830371  0.7347193  0.67728899]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0063.jpg:  
Plane Normal: [-0.03621323 -0.40195618  0.91494253]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0045.jpg:  
Plane Normal: [-0.03928982 -0.09497599  0.99470391]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0044.jpg:  
Plane Normal: [ 0.01879283 -0.84186617  0.99894647]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0065.jpg:  
Plane Normal: [ 0.0929316 -0.0884278  0.99173799]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0072.jpg:  
Plane Normal: [ 0.05979655 -0.09009025  0.99413687]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0069.jpg:  
Plane Normal: [-0.16848472 -0.1332989  0.97664953]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0041.jpg:  
Plane Normal: [-0.08240916  0.31712588  0.94479623]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0051.jpg:  
Plane Normal: [ 0.00981994 -0.17867083  0.9838599 ]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0048.jpg:  
Plane Normal: [ 0.18864601 -0.03126409  0.98154737]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0058.jpg:  
Plane Normal: [-0.02272851 -0.11248126  0.99339387]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0049.jpg:  
Plane Normal: [-0.11487977 -0.42405507  0.89832862]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0064.jpg:  
Plane Normal: [ 0.07937017 -0.02932032  0.99641392]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0062.jpg:  
Plane Normal: [-0.09642944  0.38968739  0.91588487]  
/kaggle/input/chess-board/Dataset2/Dataset2/IMG-20250404-WA0059.jpg:  
Plane Normal: [-0.03350284  0.35900989  0.93273228]
```

Panorama Generation

Cluster 1

Image 1



Image 2



Image 3



Image 4

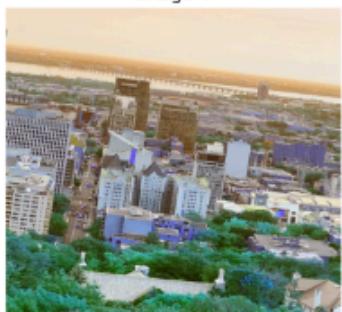


Image 5



Image 6



Image 7

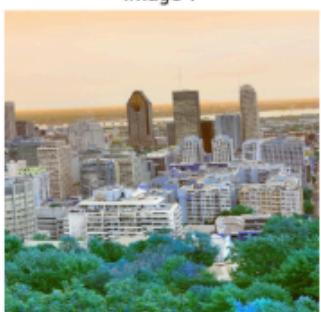
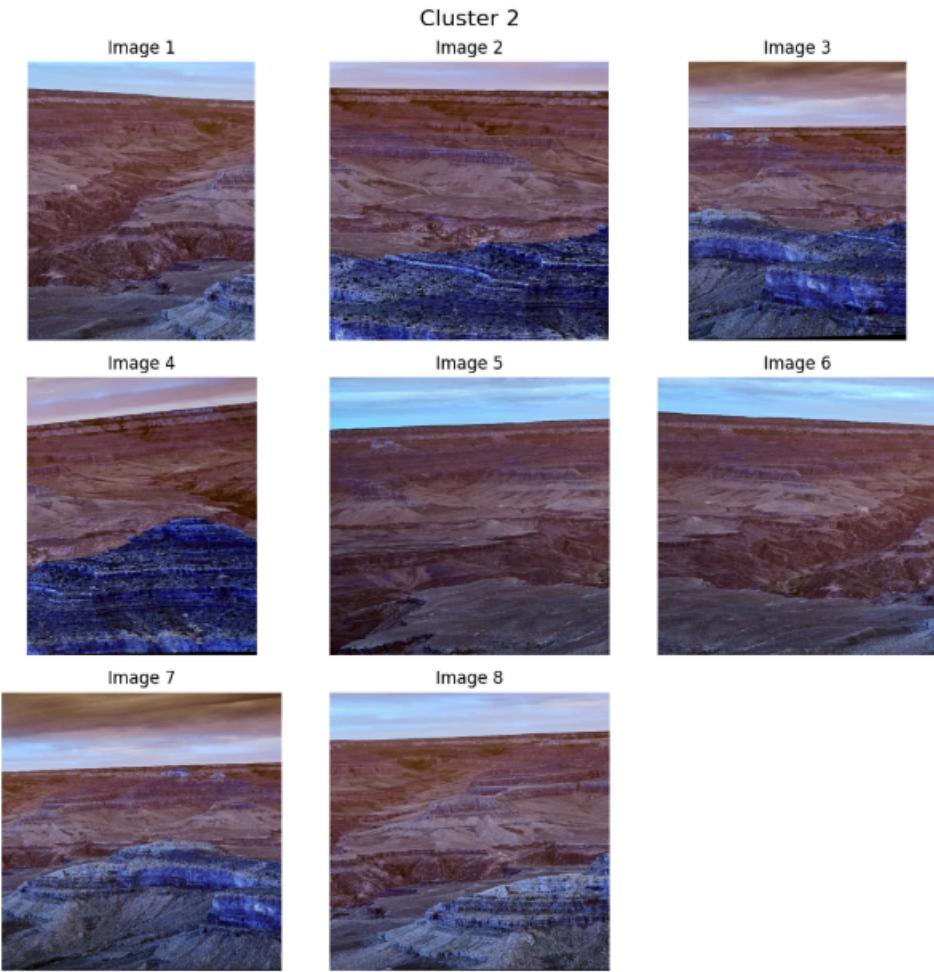


Image 8





Cluster 3



Using K means

Cluster 1

Image 1



Image 2



Cluster 2

Image 1



Image 2

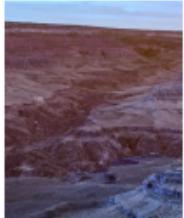


Image 3

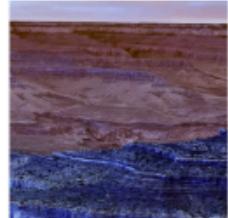


Image 4



Image 5



Image 6



Image 7



Image 8



Image 9



Image 10



Image 11

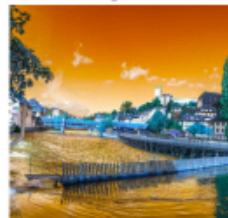


Image 12



Image 13

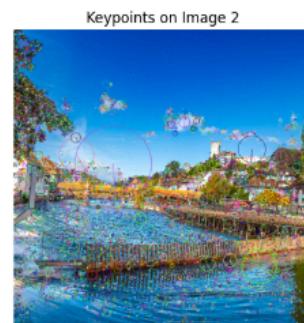
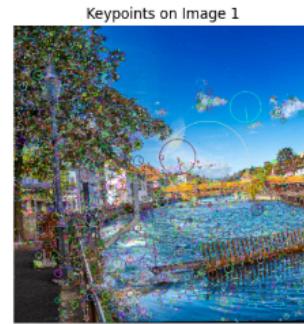




Using VBOW

K-means clustering based on their color histograms performs better than Visual Bag of Words.

1)

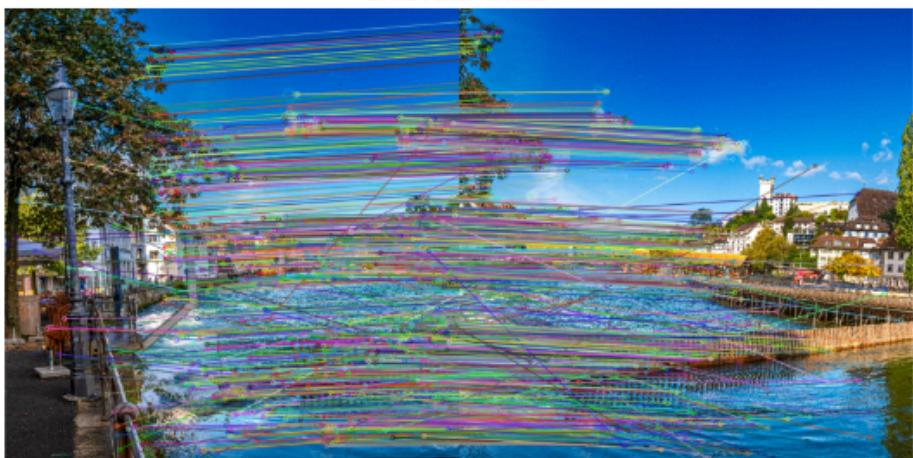


```
23]:  
bf_matcher = cv2.BFMatcher(cv2.NORM_L2, crossCheck=True)
```

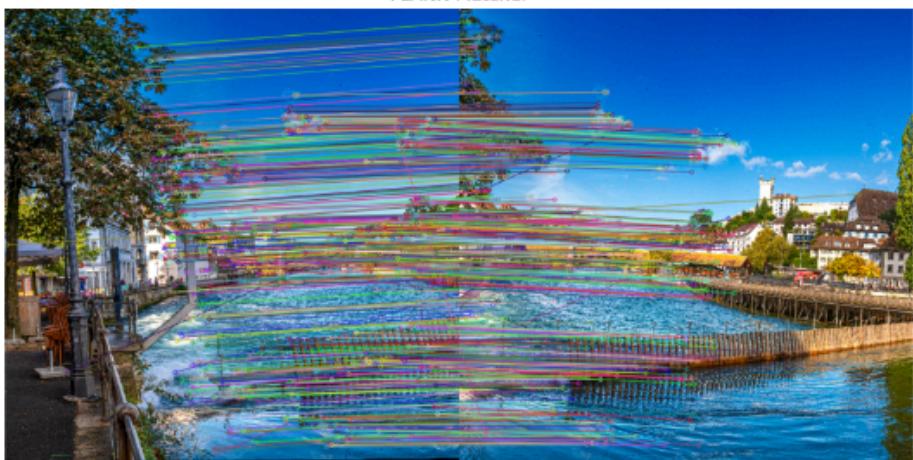
The extracted keypoints and descriptors are considered correct because they are accurately localized on distinct and meaningful features within the original images, such as corners, edges, and textured regions. The keypoints are well-distributed across the image, ensuring comprehensive feature representation, and they align consistently with the structural details of the scene. The descriptors associated with these keypoints show strong consistency, indicating reliable feature extraction. Additionally, when matched with keypoints from other images, they produce correct correspondences, further confirming their accuracy. This alignment and consistency demonstrate that the keypoints and descriptors are correctly identified and suitable for subsequent image processing tasks.

2)

Brute Force Matcher



FLANN Matcher



3)

```
[[ 8.98085298e-01 -1.32630096e-01 -1.63930197e+02]
 [ 1.43784982e-01  9.88078715e-01 -6.93765571e+01]
 [ 1.50739776e-06 -6.92855471e-06  1.00000000e+00]]
Homography matrix saved as 'homography_matrix.csv'
```

4)

Warped Image 1



Original Image 2

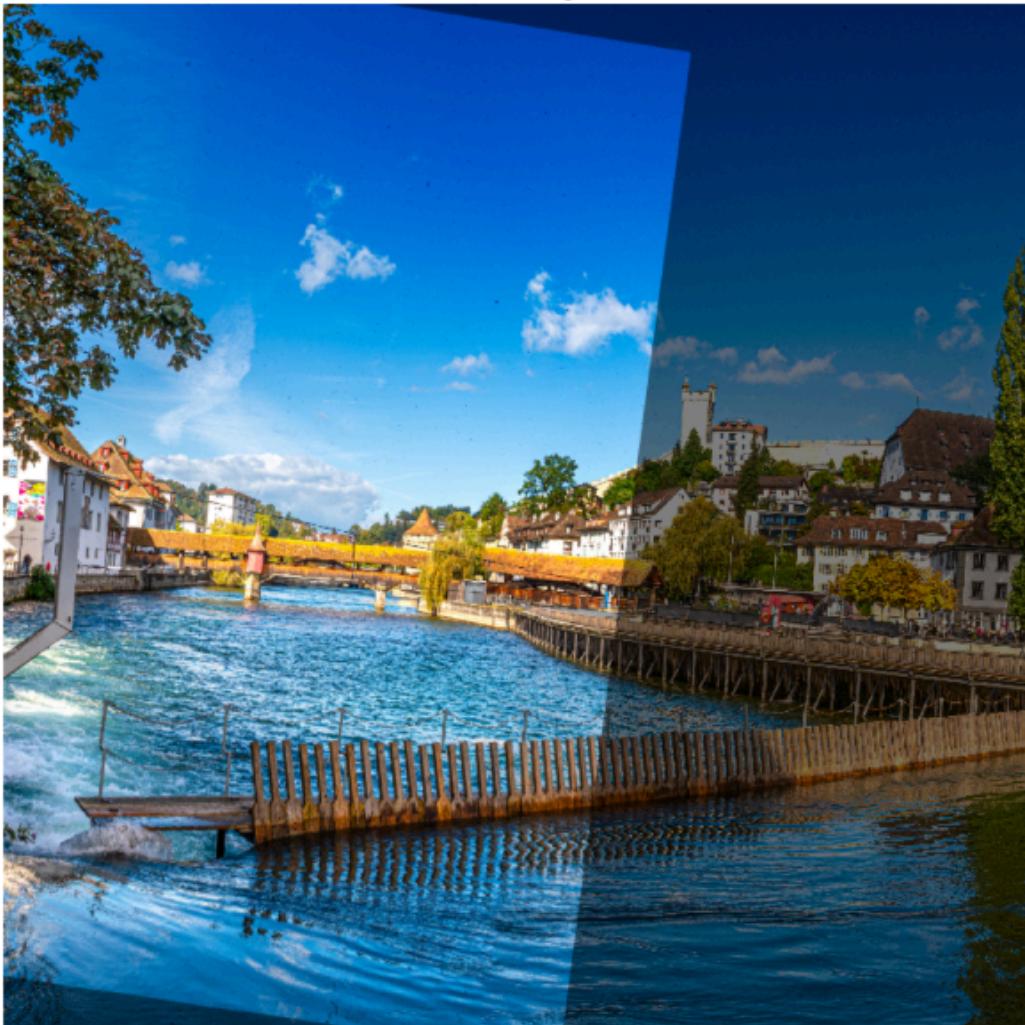


7]: (-0.5, 1439.5, 719.5, -0.5)

Panorama without Cropping and Blending



Blended Image



6)

Stitched Panorama (8 images)



Stitched Panorama (8 images)



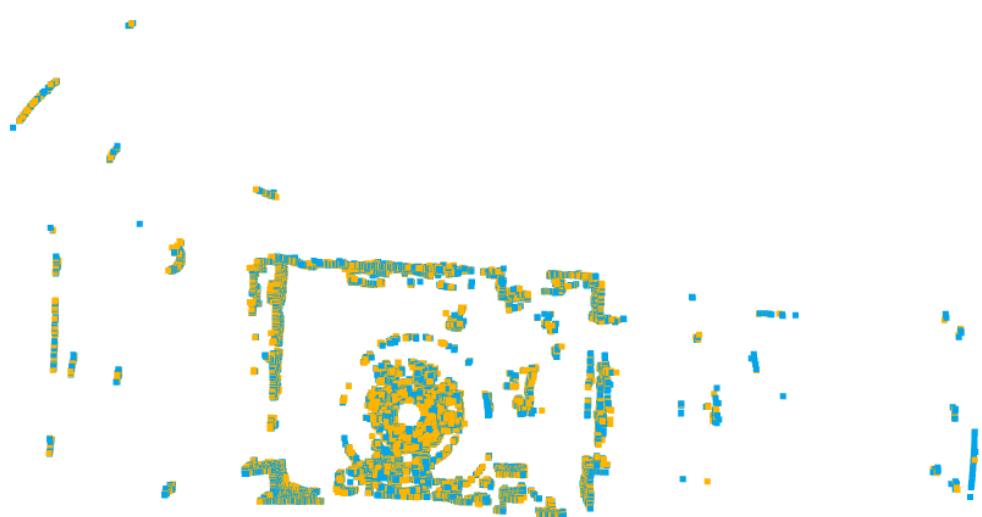
Panorama saved as 'panorama_result.jpg'



Note: I wrote the code of stitching in this part it was not performing well so i took that code to deepseek for debugging made few changes as suggested by deepseek

BONUS

1)



```

Initial alignment: Fitness 0.9873, Inlier RMSE 0.0198
Initial Transformation Matrix:
[[1. 0. 0. 0.]
 [0. 1. 0. 0.]
 [0. 0. 1. 0.]
 [0. 0. 0. 1.]]
ICP Transformation Matrix:
[[ 9.9999988e-01 -1.48888745e-04  3.74253766e-05 -4.38870580e-03]
 [ 1.48888602e-04  9.99999989e-01  3.81676438e-06  2.58107696e-05]
 [-3.74259444e-05 -3.81119213e-06  9.99999999e-01  1.09288603e-06]
 [ 0.00000000e+00  0.00000000e+00  0.00000000e+00  1.00000000e+00]]
ICP alignment: Fitness 0.9871, Inlier RMSE 0.0194

```

2)

Experiment Results:						
	Initial Guess	Voxel Size	Threshold	Fitness	RMSE	\
0	Identity Matrix	0.05	0.1	0.993744	0.020600	
1	Random Orthogonal Matrix	0.05	0.1	0.044308	0.063899	
2	RANSAC Initial Guess	0.05	0.1	0.993744	0.020600	
3	Identity Matrix	0.05	0.2	0.999077	0.022497	
4	Random Orthogonal Matrix	0.05	0.2	0.177538	0.111721	
5	RANSAC Initial Guess	0.05	0.2	0.999077	0.022497	
6	Identity Matrix	0.05	0.3	0.999795	0.023471	
7	Random Orthogonal Matrix	0.05	0.3	0.382154	0.176361	
8	RANSAC Initial Guess	0.05	0.3	0.999795	0.023471	
9	Identity Matrix	0.10	0.1	0.987469	0.025173	
10	Random Orthogonal Matrix	0.10	0.1	0.063414	0.068491	
11	RANSAC Initial Guess	0.10	0.1	0.987469	0.025173	
12	Identity Matrix	0.10	0.2	0.998101	0.028278	
13	Random Orthogonal Matrix	0.10	0.2	0.220999	0.127462	
14	RANSAC Initial Guess	0.10	0.2	0.998101	0.028278	
15	Identity Matrix	0.10	0.3	0.999620	0.029846	
16	Random Orthogonal Matrix	0.10	0.3	0.322575	0.161652	
17	RANSAC Initial Guess	0.10	0.3	0.999620	0.029846	
18	Identity Matrix	0.20	0.1	0.969792	0.027904	
19	Random Orthogonal Matrix	0.20	0.1	0.013731	0.068064	
20	RANSAC Initial Guess	0.20	0.1	0.969792	0.027904	
21	Identity Matrix	0.20	0.2	0.996469	0.035424	
22	Random Orthogonal Matrix	0.20	0.2	0.045116	0.123717	
23	RANSAC Initial Guess	0.20	0.2	0.996469	0.035424	
24	Identity Matrix	0.20	0.3	0.999215	0.038014	
25	Random Orthogonal Matrix	0.20	0.3	0.083170	0.172619	
26	RANSAC Initial Guess	0.20	0.3	0.999215	0.038014	

```
Best Parameters based on Fitness:  
Initial Guess                                Identity Matrix  
Voxel Size                                     0.05  
Threshold                                      0.3  
Fitness                                         0.999795  
RMSE                                            0.023471  
Transformation Matrix   [[0.999999841837576, -0.000171576610114649, 4...  
Name: 6, dtype: object  
  
Best Parameters based on RMSE:  
Initial Guess                                RANSAC Initial Guess  
Voxel Size                                     0.05  
Threshold                                      0.1  
Fitness                                         0.993744  
RMSE                                            0.0206
```

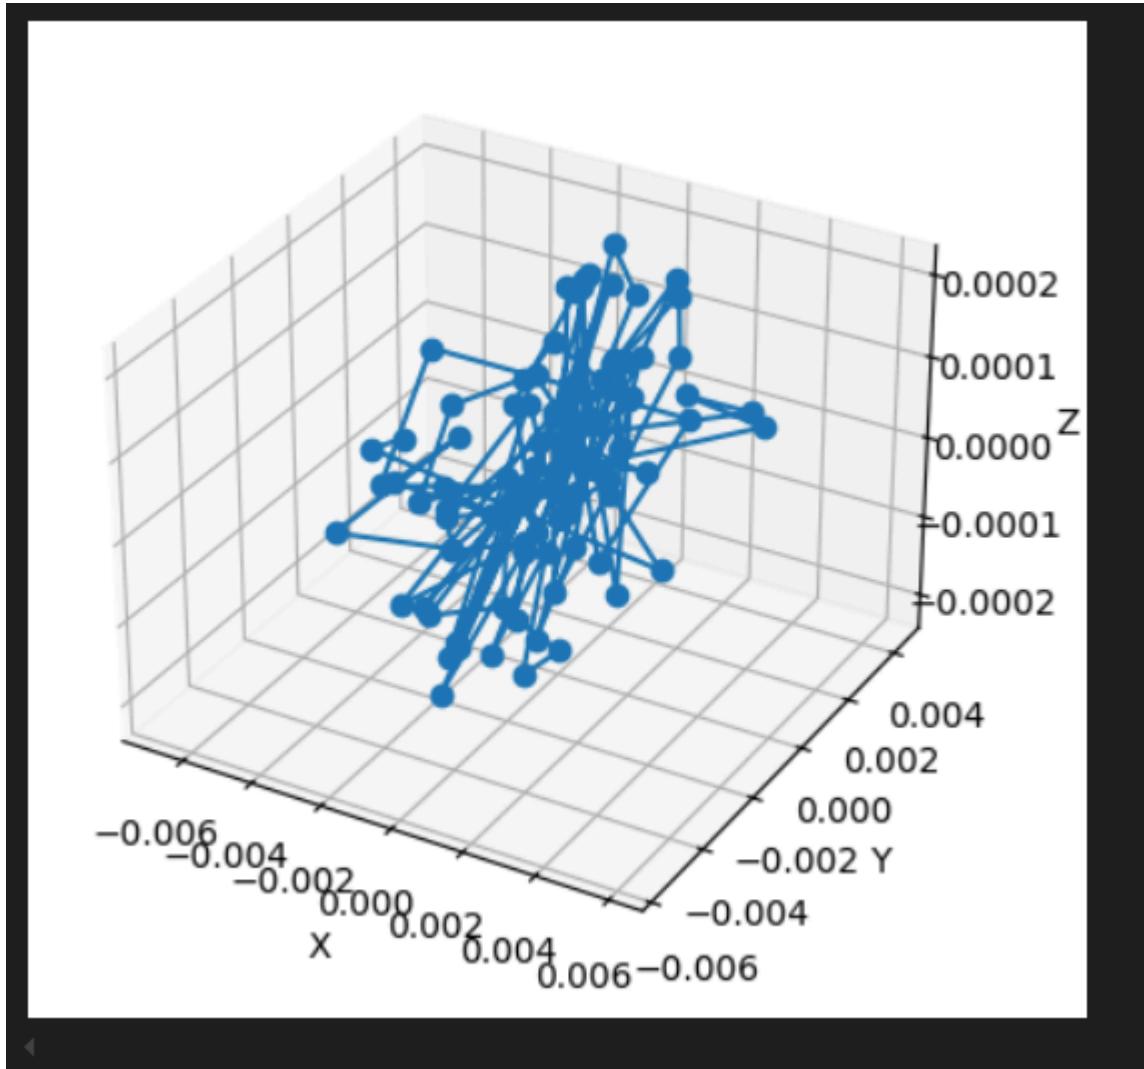
```
array([[ 9.9999984e-01, -1.71576610e-04,  4.68396288e-05,  
       -4.28356978e-03],  
      [ 1.71576591e-04,  9.9999985e-01,  4.11195508e-07,  
       -4.36617041e-04],  
      [-4.68396986e-05, -4.03158918e-07,  9.9999999e-01,  
       3.00209219e-05],  
      [ 0.00000000e+00,  0.00000000e+00,  0.00000000e+00,  
       1.00000000e+00]])
```



Observations:

- **Good alignment** was achieved with minimal drift.
- **Some minor misalignment** was still present in areas with fewer features.

4)



For Whole Question I took help of deep seek for debugging purposes

References

- 1)<https://www.deepseek.com/>
- 2)https://docs.opencv.org/4.x/dc/dbb/tutorial_py_calibration.html
- 3)<https://theailearner.com/tag/cv2-warpperspective/>
- 4)<https://learnopencv.com/tag/ransac/>
- 5)https://docs.opencv.org/3.4/d1/de0/tutorial_py_feature_homography.html
- 6)https://docs.opencv.org/3.4/dc/de2/classcv_1_1FlannBasedMatcher.html
- 7)https://docs.opencv.org/4.x/da/df5/tutorial_py_sift_intro.html
- 8)https://www.open3d.org/docs/latest/tutorial/pipelines/icp_registration.html